

CPU Scheduling: CFS and EEVDF


By: James Young





CPU Scheduling


- CPU schedulers allocate CPU time to various processes and threads
- Main goals are to optimize system performance, ensure fairness, and minimize latency
- Two prominent Linux CPU scheduling algorithms include Completely Fair Scheduler (CFS) and Earliest Eligible Virtual Deadline First (EEVDF)





CFS Overview

- Before CFS, Linux used CPU schedulers such as $O(n)$ scheduler and $O(1)$ scheduler
- Problem: Older CPU schedulers had problem of being either inefficient, having high complexity, or lacked fairness
- Solution: CFS introduced into Linux in kernel 2.6, where it focuses on fairness whilst having efficient operations






CFS Algorithm Implementation

- CFS algorithm uses **virtual runtime** (vruntime) and **red-black tree**
- For fairness, CFS computes vruntime for each process and chooses the smallest vruntime to run
- For efficiency, CFS stores the processes in a red-black tree, where each node is a process with a value vruntime

$$vruntime = (actual\ runtime) \cdot \frac{1024}{1 + nice\ value}$$




Example of a CFS Red-Black Tree



CFS Red-Black Tree

- Red-Black Tree Properties:
 - a. Root Property: The root is always black.
 - b. Red Property: Red nodes cannot have red children, preventing consecutive red nodes.
 - c. Black Depth Property: Every path from a node to its descendant leaves must have the same number of black nodes.
 - d. Leaf Nodes: All null leaf nodes are black.
- Above properties keep tree balanced, meaning efficient operations ($O(\log n)$ time complexity)
- Red-Black tree structure also ensures that leftmost node will be smallest virtual runtime value, so CFS chooses leftmost node to run



Example of a Red-Black Tree



CFS Example

Store 4 Processes (Figure 1):

- Process A: vruntime = 1
- Process B: vruntime = 3
- Process C: vruntime = 5
- Process D: vruntime = 7




Figure 1

Insert Process E (Figure 2):

- Process E: vruntime = 2




Figure 2

Run Process (Figure 3):

- CFS chooses lowest vruntime, so leftmost node (Process A)
- After deletion of node, tree rebalances





Figure 3



EEVDF Overview

- CFS Problem: Does not explicitly consider latency requirements for processes
- Relies instead on nice values for priority, which does not adequately address varying latency needs
- Solution: EEVDF, which uses concepts of **lag** and **virtual deadline** to remain fair whilst minimizing latency, replaced CFS in kernel 6.6





EEVDF Algorithm: Lag

- To ensure fairness, EEVDF calculates lag to determine a process eligibility
- Lag is basically the difference between the time that process should have gotten and how much it actually got
- If a process has positive or zero lag value, it is eligible to be run and negative lag value means it is not eligible
- Tasks with high lag are prioritized, aim to even out lag values across and avoid excessive delays



$$lag = Expected\ Execution\ Time - Actual\ Execution\ Time$$



EEVDF Algorithm: Virtual Deadline

- To minimize latency, EEVDF calculates virtual deadline, which takes into account a process' urgency and expected execution time
- Ensures that processes with the highest need for CPU time are scheduled promptly, quicker access for latency-sensitive processes
- Done by choosing the smallest virtual deadline (Earliest Deadline) task among the Eligible Tasks
- Stores processes in Red-Black tree (similar to CFS), with virtual deadline being the node value











$$virtual\ deadline = eligible\ time + \left(\frac{time\ slice}{weight} \right)$$



CFS vs EEVDF Performance Benchmark

- Ran benchmarks to compare performance, used Phoronix Test Suite for running benchmarks
- Scheduler-based benchmark: Hackbench (benchmark tool for kernel schedule)
- Real-world Application benchmark: Nginx, 7-zip, C-Ray, FFmpeg







References

- N. Goel and R. B. Garg, A Comparative Study of CPU Scheduling Algorithms. 2013. [Online]. Available: <https://arxiv.org/abs/1307.4165>
- M. Jones, “Inside the Linux 2.6 Completely Fair Scheduler,” IBM developer, <https://developer.ibm.com/tutorials/l-completely-fair-scheduler/> (accessed Nov. 5, 2024).
- P. Patil, S. Dhotre, and R. Jamale, “A Survey on Fairness and Performance Analysis of Completely Fair Scheduler in Linux Kernel,” International Journal of Control Theory and Applications, vol. 9, no. 44, 2016.
- J. Corbet, “An EEVDF CPU scheduler for linux,” LWN.net, <https://lwn.net/Articles/925371/> (accessed Nov. 8, 2024).
- Stoica, I., & Abdel-WahabDepartment, H. (1995). Eligible Virtual Deadline First : A Flexibleand Accurate Mechanism for Proportional ShareResource Allocation. <https://api.semanticscholar.org/CorpusID:7263527>
- “CFS Scheduler,” The Linux Kernel Documentation, <https://docs.kernel.org/scheduler/sched-design-CFS.html> (accessed Nov. 8, 2024).
- “EEVDF Scheduler,” The Linux Kernel Documentation, <https://docs.kernel.org/scheduler/sched-eevdf.html> (accessed Nov. 8, 2024).