

ELEC 3210

Introduction to Mobile Robotics

Lecture 15

(Machine Learning and Information Processing for Robotics)

Huan YIN

Research Assistant Professor, Dept. of ECE

eehyin@ust.hk




Recap L14

- Harris Corner
 - Eigen value and eigen vectors



Feature Descriptors

- We know how to detect good points
- Next question: **How to match them?**



Feature Descriptor and Matching

Invariance

- Suppose we are comparing two images I_1 and I_2
 - I_2 may be a transformed version of I_1
 - What kinds of transformations are we likely to encounter in practice?
- We'd like to find the same features regardless of the transformation
 - This is called transformational ***invariance***
 - Most feature methods are designed to be invariant to
 - Translation, 2D rotation, scale
 - They can usually also handle
 - Limited 3D rotations (SIFT works up to about 60 degrees)
 - Limited illumination/contrast changes

How to achieve invariance

Need both of the following:

1. Make sure your detector is invariant

- Harris is invariant to translation and rotation
- Scale is trickier
 - Many sophisticated methods find “the best scale” to represent each feature (e.g., SIFT), Autoscale in Lecture 14

2. Design an invariant feature *descriptor*

- A descriptor captures the information in a region around the detected feature point
- The simplest descriptor: a square window of pixels
 - What’s this invariant to?
- Let’s look at some better approaches...

Scale Invariant Feature Transform

- Basic idea:
 - Take 16x16 square window around detected feature
 - Compute edge orientation for each pixel
 - Throw out weak edges (threshold gradient magnitude)
 - Create histogram of surviving edge orientations







Image gradients




- Full version
 - Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
 - Compute an orientation histogram for each cell
 - 16 cells * 8 orientations = 128 dimensional descriptor



SIFT Properties

- Translation invariance
- Illumination invariance
- Rotation invariance
 - rotate all pixel gradients to 90 deg using to the “main direction”
 - main direction = the highest bin in histogram
 - invariant up to about 60 degree out of plane rotation




Problem

Given a feature in I_1 , how to find the best match in I_2 ?


1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min distance

Feature Distance

- How to define the difference between two features f_1, f_2 ?
 - Simple approach is $SSD(f_1, f_2)$
 - sum of square differences between entries of the two descriptors
 - can give good scores to very ambiguous (bad) matches




I_1




I_2

Feature Distance

- How to define the difference between two features f_1, f_2 ?
 - Better approach: ratio distance = $\text{SSD}(f_1, f_2) / \text{SSD}(f_1, f_2')$
 - f_2 is best SSD match to f_1 in I_2
 - f_2' is 2nd best SSD match to f_1 in I_2
 - gives small values for ambiguous matches

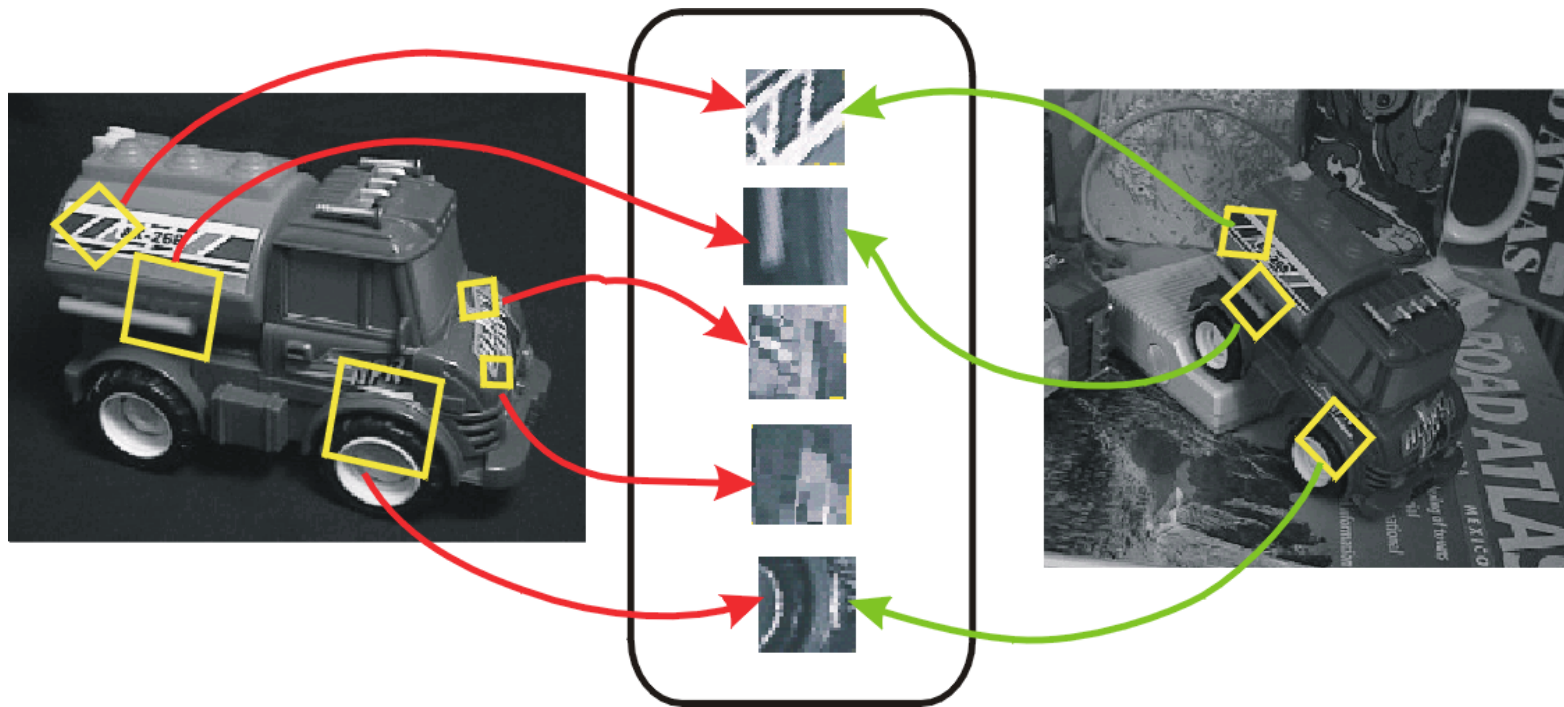


I_1




I_2

Feature Matching by SIFT



3D-2D Pose Estimation (Pose from Projective Transform) (The PnP Problem)


Recap L4 - Depth from Stereo Vision



Recap L5 ICP



- Pose estimation via Iterative Closest Point
- 3D-3D pose estimation

Iteration 0





3D-3D for Visual Pose Estimation

- How to obtain 3D-3D data association?
 - Calibrated stereo image pairs as input
 - Spatial matching – feature matching between stereo image pairs, for computation of 3D points
 - Temporal matching – feature matching between images captures at different times, for motion estimation
 - Need to address outlier removal (RANSAC)– to be discussed soon
 - Usually poor performance due to increased ranging error at longer distance with stereo vision – use **3D-2D pose estimation** instead




How about 2D-3D?



- 3D features in frame 1
(or known 3D features in world frame)
- △ 2D feature observations in frame 2
(or feature observations in body frame)
- 2D feature reprojections of given a
estimated pose of frame 2 in frame 1
(or estimated transformation of body
frame in world frame)

3D-2D

Recap L4 Pin-hole Camera Model

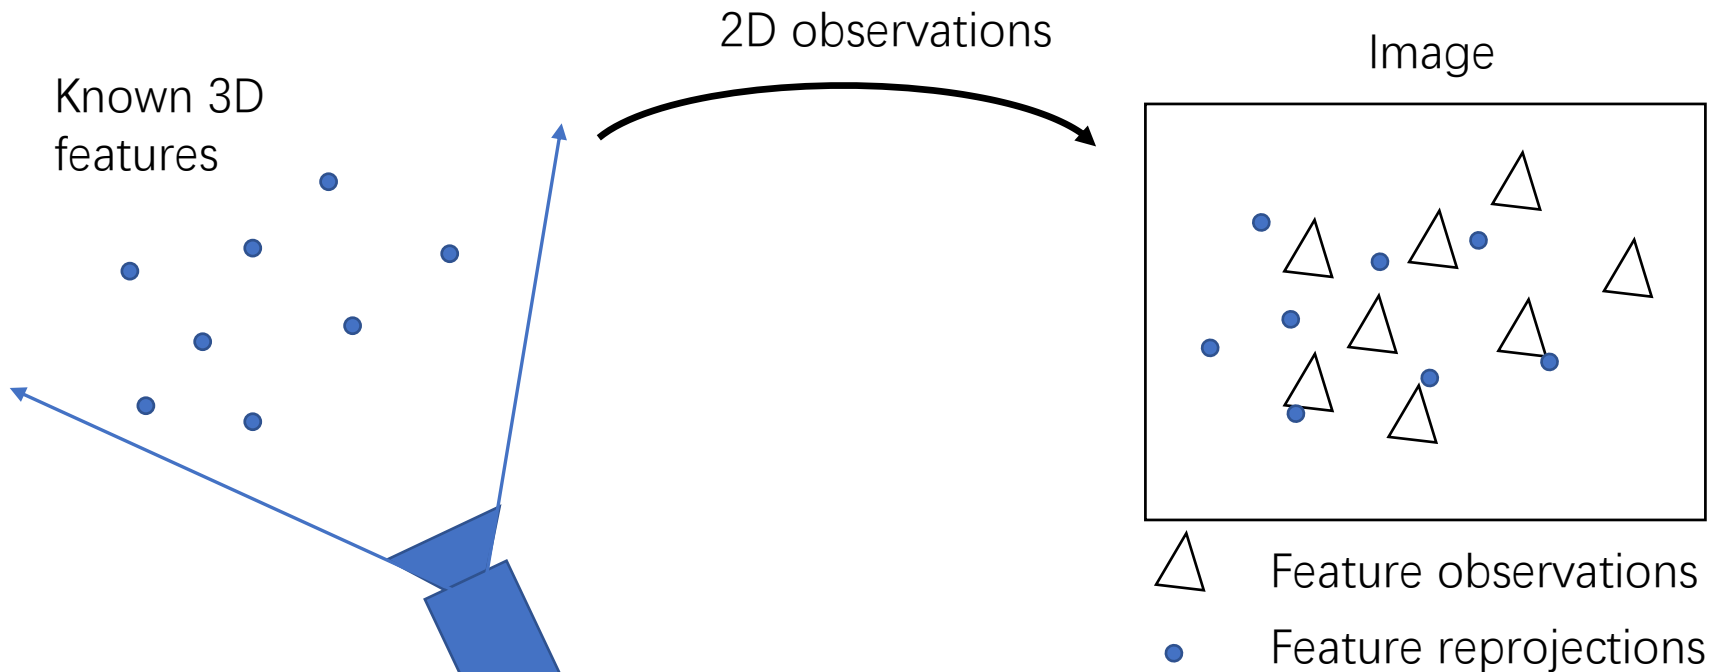


Nonlinear 3D-2D Pose Estimation


θ : Euler Angles $\in R^3$ t : Translation $\in R^3$ $\pi(\cdot)$: projection function

- Minimize the reprojection error w.r.t. camera pose
 - Can be solved via Gauss-Newton method

$$\min_{\theta, t} \sum_i \left\| \begin{bmatrix} u_i \\ v_i \end{bmatrix} - \pi \left(K \cdot (R(\theta) \cdot \begin{bmatrix} x_i \\ y_i \\ z_i \end{bmatrix} + t) \right) \right\|^2$$




3D-2D Visual Localization with Markers



Courtesy: Shaojie Shen

What if no markers?



- In unstructured environments
 - For data association



$$\min_{\theta, t} \sum_i \left\| \begin{bmatrix} u_i \\ v_i \end{bmatrix} - \pi \left(\overset{???}{K} \cdot (R(\theta) \cdot \begin{bmatrix} x_i \\ y_i \\ z_i \end{bmatrix} + t) \right) \right\|^2$$

No Markers

- What if you do not have markers?
 - Outlier rejection




RANdom SAmple Consensus (RANSAC)

- Model fitting with outlier rejection
 - The 6-DOF pose you are trying to estimation is a model
- Algorithm:
 - Loop:
 - Randomly select a small amount of (or minimum) data points to find a model
 - See the error between the model and all other data points
 - Find the data points with error smaller than a threshold as inliers
 - If the current model has more inliers than all previous ones, record all inliers
 - Repeat
 - Use all inliers to find the best estimate of the model

RANSAC


- RANSAC for 2D line fitting
 - Minimum number of points to define a 2D line: 2
 - Error metric: point to line distance
 - How many iterations are required?

- Iteration 1: 4 Inliers




RANSAC

- RANSAC for 2D line fitting
 - Minimum number of points to define a 2D line: 2
 - Error metric: point to line distance
 - How many iterations are required?
- Iteration 2: 10 Inliers




- RANSAC for 2D line fitting
 - Minimum number of points to define a 2D line: 2
 - Error metric: point to line distance
 - How many iterations are required?
- Iteration 3: 12 Inliers



Failure

- RANSAC for 2D line fitting
 - Minimum number of points to define a 2D line: 2
 - Error metric: point to line distance
 - How many iterations are required?
- Iteration 4: 4 Inliers



Pose Estimation with RANSAC

- How many feature correspondences are required to create a model?
 - 3D-3D: 3
 - 3D-2D: 3
 - It is OK to use more points to find the model
 - But few number of points is better (Why?)
- How to define the error metric?
 - 3D-3D: point distance
 - 3D-2D: reprojection error
- How many iterations are required?


- How many iterations are required - the probability
 - Probability of outlier: X ($X < 1$)
 - M number of data points to create a model
 - N iterations
- RANSAC failure: all random sample contains at least 1 outlier
 - Failure probability = $(1 - (1 - X)^M)^N$
 - 2D line fitting example
 - 30% outliers
 - 2 data points to create a model
 - Failure probability for 5 iterations: 3.45%
 - Failure probability for 10 iterations: 0.12%

- How many iterations are required - the probability
 - Probability of outlier: X ($X < 1$)
 - M number of data points to create a model
 - N iterations
- RANSAC failure: all random sample contains at least 1 outlier
 - Failure probability = $(1 - (1 - X)^M)^N$
 - 3D-3D pose estimation
 - 30% outliers
 - 3 data points to create a model
 - Failure probability for 5 iterations: 12.24%
 - Failure probability for 10 iterations: 1.49%
 - Failure probability for 20 iterations: 0.02%

- How many iterations are required - the probability
 - Probability of outlier: X ($X < 1$)
 - M number of data points to create a model
 - N iterations
- RANSAC failure: all random sample contains at least 1 outlier
 - Failure probability = $(1 - (1 - X)^M)^N$
 - 3D-3D pose estimation
 - 30% outliers
 - 20 data points to create a model (bad example)
 - Failure probability for 5 iterations: 99.6%
 - Failure probability for 10 iterations: 99.2%
 - Failure probability for 20 iterations: 98.4%
 - Failure probability for 1000 iterations: 45%
 - Failure probability for 10000 iterations: 0.03%


RANSAC for Alignment

- Extract features




RANSAC for Alignment

- Extract features
- Compute putative matches




RANSAC for Alignment

- Extract features
- Compute putative matches
- Loop:
 - Hypothesize transformation T
 - Verify transformation (search for other matches consistent with T)



RANSAC for Alignment


- Extract features
- Compute putative matches
- Loop:
 - Hypothesize transformation T
 - Verify transformation (search for other matches consistent with T)



Visual Odometry

Pipeline

- A mini pipeline, but can work



Visual Odometry



```
frame          1
key            0
keyframes      1
from start     0.001m
covered        0.000m
inliers        319
outliers       10
time per frame 34ms
```

```
FeatureDetectorFast
DescriptorSchemeSAD
```


Deep Learning-based

- Sarlin PE, DeTone D, Malisiewicz T, Rabinovich A. Superglue: Learning feature matching with graph neural networks. CVPR 2020 (pp. 4938-4947).



SuperGlue: Learning Feature Matching with Graph Neural Networks

Paul-Edouard Sarlin¹
Tomasz Malisiewicz²

Daniel DeTone²
Andrew Rabinovich²



Next Lecture

- Sensing + Estimation 😊
- Planning

