

ELEC 3210

Introduction to Mobile Robotics

Lecture 17

(Machine Learning and Information Processing for Robotics)

Huan YIN


Research Assistant Professor, Dept. of ECE

eehyin@ust.hk



Recap L16

R.O.B.O.T. Comics



"HIS PATH-PLANNING MAY BE
SUB-OPTIMAL, BUT IT'S GOT FLAIR."

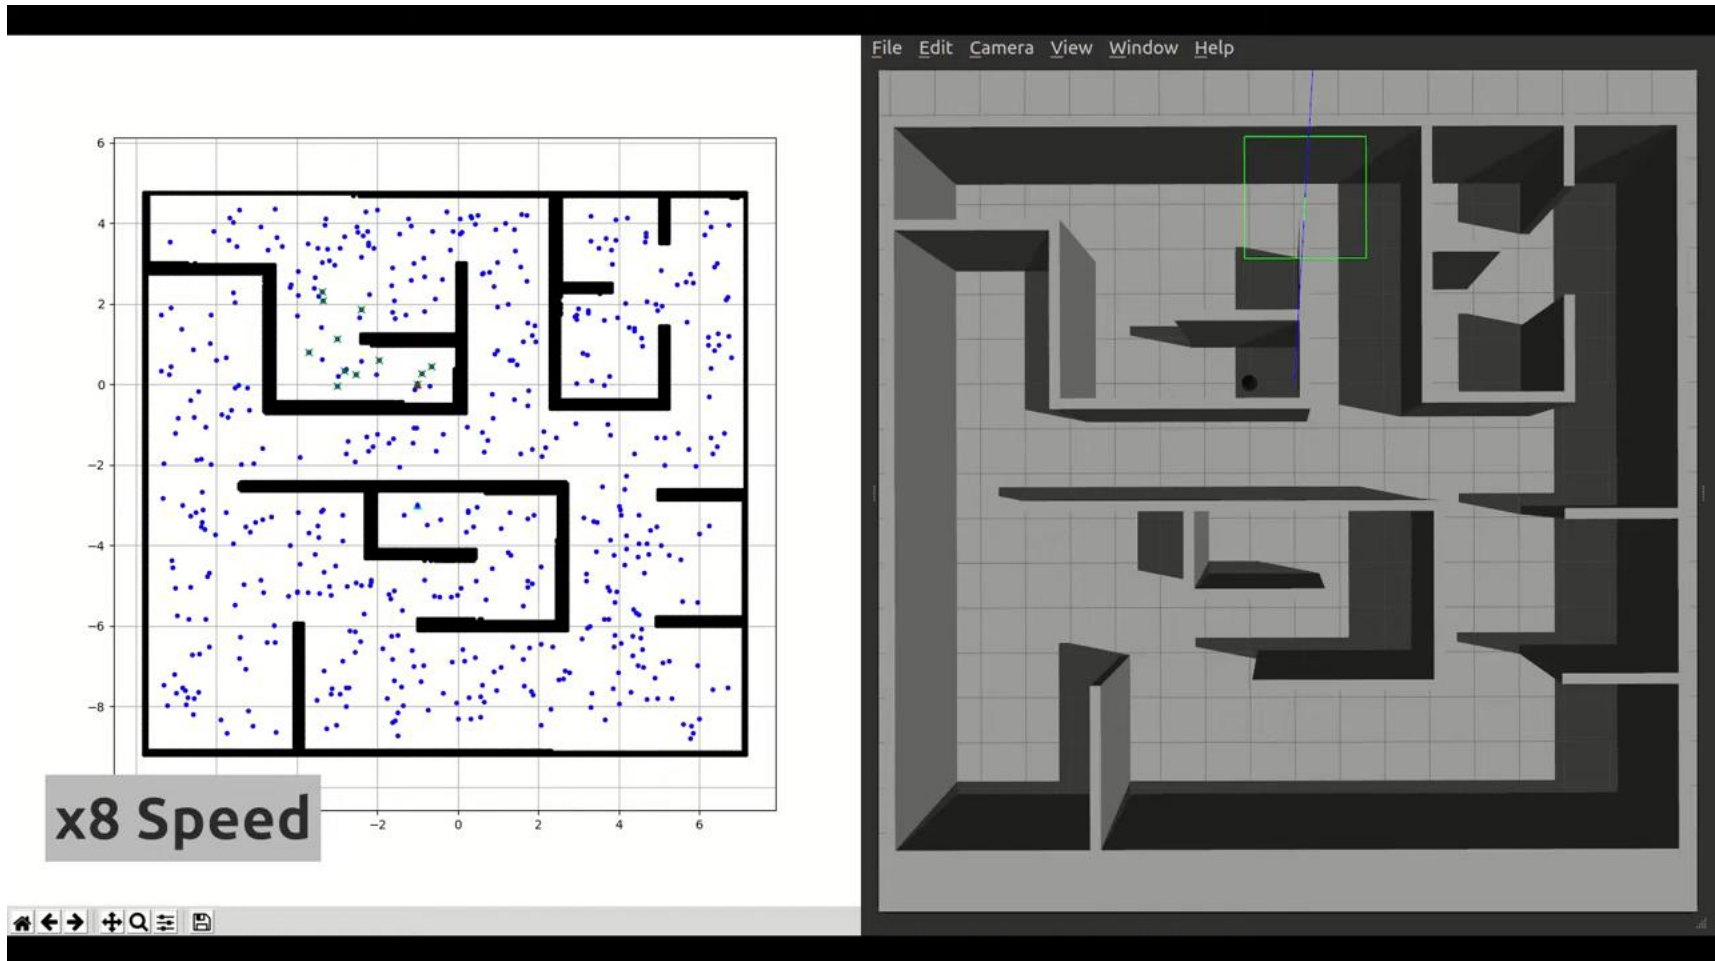
Recap L16

- Basic concepts of Motion Planning
- Planning as graph search problem
- How to construct the graph?

- Combinatorial Planning
 - Resolution Completeness
 - Visibility Graph, Voronoi Diagram, Cell Decomposition
- Sampling-based Planning
 - Probabilistic Completeness
 - Probabilistic road maps (PRM)
 - Rapidly exploring random tree (RRT)


RRT

Probabilistic road maps (PRM)




Probabilistically Complete

- Do not work well for some problems, narrow passages




Rapidly Exploring Random Trees

- **Idea:** aggressively probe and explore the C-space by expanding incrementally from an initial configuration
- The explored territory is marked by a tree rooted at the initial



45 iterations



2345 iterations

- The algorithm: Given \mathcal{C} and q_0

Algorithm 1: RRT

```
1  $G.\text{init}(q_0)$ 
2 repeat
3    $q_{\text{rand}} \rightarrow \text{RANDOM\_CONFIG}(\mathcal{C})$ 
4    $q_{\text{near}} \leftarrow \text{NEAREST}(G, q_{\text{rand}})$ 
5    $G.\text{add\_edge}(q_{\text{near}}, q_{\text{rand}})$ 
6 until condition
```



Sample from a bounded region
centered around q_0



- The algorithm: Given \mathcal{C} and q_0

Algorithm 1: RRT

```
1  $G.\text{init}(q_0)$ 
2 repeat
3    $q_{\text{rand}} \rightarrow \text{RANDOM\_CONFIG}(\mathcal{C})$ 
4    $q_{\text{near}} \leftarrow \text{NEAREST}(G, q_{\text{rand}})$ 
5    $G.\text{add\_edge}(q_{\text{near}}, q_{\text{rand}})$ 
6 until condition
```



Finds closest vertex in G using a distance function



- The algorithm: Given \mathcal{C} and q_0


Algorithm 1: RRT

```
1  $G.\text{init}(q_0)$ 
2 repeat
3    $q_{\text{rand}} \rightarrow \text{RANDOM\_CONFIG}(\mathcal{C})$ 
4    $q_{\text{near}} \leftarrow \text{NEAREST}(G, q_{\text{rand}})$ 
5    $G.\text{add\_edge}(q_{\text{near}}, q_{\text{rand}})$ 
6 until condition
```



Several strategies to find q_{near} given the closest vertex on G :

- take closest vertex
- Check intermediate points at regular intervals and split edge at q_{near}



- The algorithm: Given \mathcal{C} and q_0


Algorithm 1: RRT

```
1  $G.\text{init}(q_0)$ 
2 repeat
3    $q_{\text{rand}} \rightarrow \text{RANDOM\_CONFIG}(\mathcal{C})$ 
4    $q_{\text{near}} \leftarrow \text{NEAREST}(G, q_{\text{rand}})$ 
5    $G.\text{add\_edge}(q_{\text{near}}, q_{\text{rand}})$ 
6 until condition
```



Connect nearest point with random point that travels from q_{near} to q_{rand}

- No collision: add edge
- Collision: new vertex is q_i , as close as possible to \mathcal{C}_{obs}



RRT

- The algorithm: Given \mathcal{C} and q_0


Algorithm 1: RRT

```
1  $G.\text{init}(q_0)$ 
2 repeat
3    $q_{\text{rand}} \rightarrow \text{RANDOM\_CONFIG}(\mathcal{C})$ 
4    $q_{\text{near}} \leftarrow \text{NEAREST}(G, q_{\text{rand}})$ 
5    $G.\text{add\_edge}(q_{\text{near}}, q_{\text{rand}})$ 
6 until condition
```




Connect nearest point with random point that travels from q_{near} to q_{rand}



- No collision: add edge
- Collision: new vertex is q_i , as close as possible to \mathcal{C}_{obs}



- RRT is exploring the space, explore until the final configuration is reached
- Can add bias to the goal when expanding randomly
- Pros:
 - Balance between greedy search and exploration
 - Easy to implement
- Cons
 - Metric sensitivity
 - Unknown rate of convergence




- Basic Idea
 - RRT is simple, but is prone to be probabilistic incomplete
 - Add **rewire** function: swap new point in as parent for nearby vertices who can be reached along shorter path through new point than through their original (current) path
 - RRT* is asymptotically optimal.




RRT*

RRT




RRT*




Source: Karaman and Frazzoli

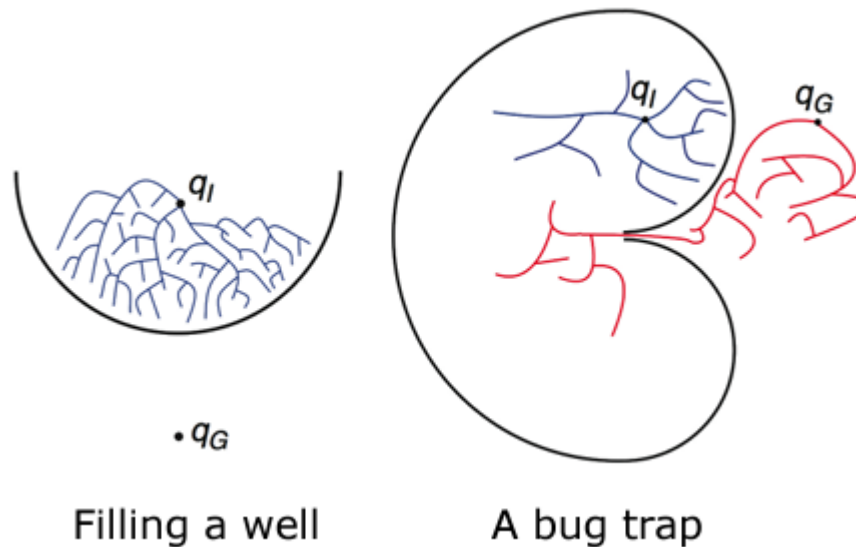
RRT



RRT*



- Some problems require more effective methods: bidirectional search
- Grow two RRTs
- In every other step, try to extend each tree towards the newest vertex of the other tree



RRT Page

- <https://lavalle.pl/rrt/>


The RRT Page




Graph Search

Search-based Method


- State space graph: a mathematical representation of a **search algorithm**
 - For every search problem, there's a corresponding state space graph
 - Connectivity between nodes in the graph is represented by (directed or undirected) edges



Grid-based graph: use grid as vertices and grid connections as edges



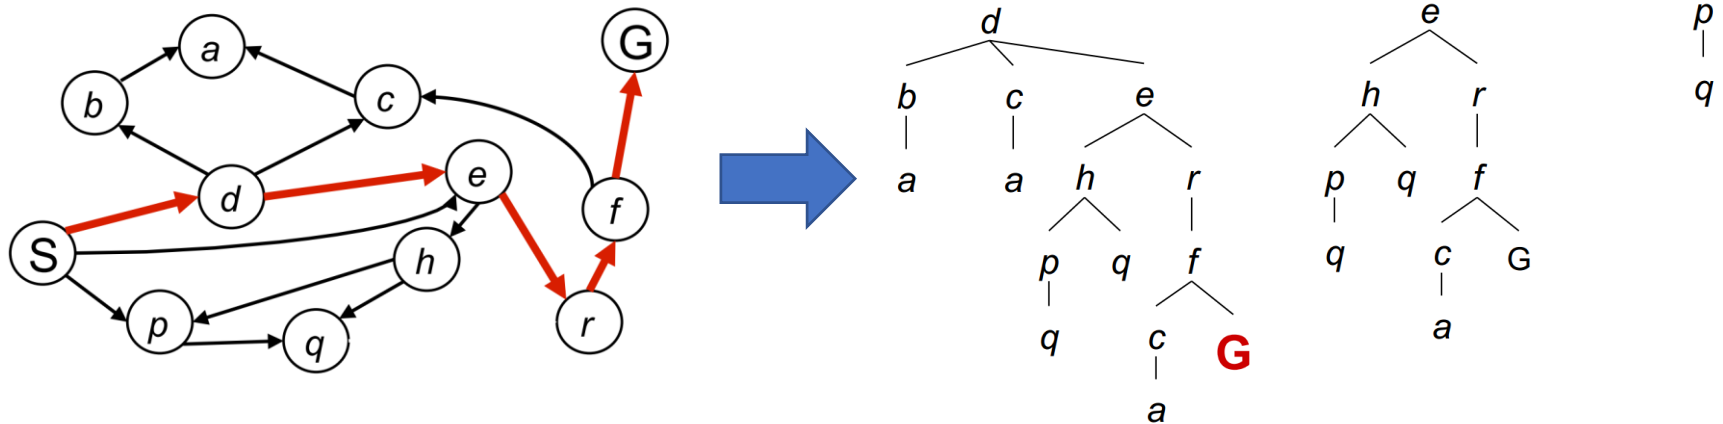
The graph generated by probabilistic roadmap (PRM)



Ridiculously tiny search graph for a tiny search problem

From Graph to Search Tree

- The search always start from start state X_s
 - Searching the graph produces a search tree, this is a “what if” tree of plans and outcomes
 - Back-tracing a node in the search tree gives us a path from the start state to that node
 - For many problems we can never actually build the whole tree, too large or inefficient – we only want to reach the goal node asap.





How to Construct a Search Tree?



- Maintain a **container** to store all the nodes **to be visited**
 - Intuition: When we “discover” a node, we store it in our “memory”. We can only visit one node at a time, but we can teleport to any node that we discover before.
- The container is initialized with the start state X_s
- Loop
 - **Remove** a node from the container according to some pre-defined score function
 - **Visit a node**
 - **Expansion**: Obtain all neighbors of the node, and push them into the container
 - Discover all its neighbors
- End Loop

Depth First Search (DFS)

- Strategy: remove (visit) the **deepest** node in the container





-





Breadth First Search (BFS)

- Strategy: remove (visit) the **shallowest** node in the container



Breadth First Search (BFS)

- Implementation: maintain a first in first out (FIFO) container (i.e. queue)




Init container

Remove node with the smallest tree level

Expansion


Add children of visited node into container

Remove node with the smallest tree level





Costs on Actions

- A practical search problem has a **cost “C”** from a node to its neighbor
 - Length, time, energy, etc.
- When all weight are 1, BFS finds the least-cost path with minimal steps
- For general cases, how to find the **least-cost path** as soon as possible?



Dijkstra's Algorithm




- Strategy: remove (visit) the node with **cheapest accumulated cost $g(n)$**
 - $g(n)$: The current best estimates of the accumulated cost from the start state to node "n"
 - Update the accumulated costs $g(m)$ for all unvisited neighbors "m" of node "n"
 - A node that has been visited is guaranteed to have the smallest cost from the start state



Dijkstra's Algorithm



- Maintain a **priority queue** to store all the nodes **to be visited**
- The priority queue is initialized with the start state X_s
- Assign $g(X_s)=0$, and $g(n)=\text{infinite}$ for all other nodes in the graph
- Loop
 - If the queue is empty, return FALSE; break;
 - **Remove** the node “n” with the lowest $g(n)$ from the priority queue
 - Mark node “n” as **visited**
 - If the node “n” is the goal state, return TRUE; break;
 - For all **unvisited** neighbors “m” of node “n”
 - If $g(m) = \text{infinite}$
 - Push node “m” into the queue
 - If $g(m) > g(n) + C_{nm}$
 - $g(m) = g(n) + C_{nm}$
 - end
- End Loop

Dijkstra's Algorithm




Dijkstra's Algorithm

- Pros
 - Complete and optimal
- Cons
 - Can only see the cost *accumulated so far* (i.e. the uniform cost), thus exploring next state in every “direction”
 - No information about goal location



Search Heuristics

- Overcome the shortcomings of uniform cost search by **inferring the least cost to goal (i.e. goal cost)**
- Designed for particular search problem
- Examples: Manhattan distance VS. Euclidean distance



A*: Combining Dijkstra's and a Heuristic



- Accumulated cost
 - $g(n)$: The current best estimates of the accumulated cost from the start state to node “n”
- Heuristic
 - $h(n)$: The **estimated least cost** from node n to goal state (i.e. goal cost)
- The least estimated cost from start state to goal state passing through node “n” is $f(n) = g(n) + h(n)$
- Strategy: remove (visit) the node with **cheapest $f(n) = g(n) + h(n)$**
 - Update the accumulated costs $g(m)$ for all unvisited neighbors “m” of node “n”
 - A node that has been visited is guaranteed to have the smallest cost from the start state


A* Algorithm

- Maintain a **priority queue** to store all the nodes **to be visited**
- The heuristic function $h(n)$ for all nodes are pre-defined
- The priority queue is initialized with the start state X_s
- Assign $g(X_s)=0$, and $g(n)=\text{infinite}$ for all other nodes in the graph
- Loop
 - If the queue is empty, return FALSE; break;
 - **Remove** the node “n” with the lowest $f(n)=g(n)+h(n)$ from the priority queue
 - Mark node “n” as **visited**
 - If the node “n” is the goal state, return TRUE; break;
 - For all **unvisited** neighbors “m” of node “n”
 - If $g(m) = \text{infinite}$
 - Push node “m” into the queue
 - If $g(m) > g(n) + C_{nm}$
 - $g(m) = g(n) + C_{nm}$
 - end
- End Loop



Only difference
comparing to Dijkstra's
algorithm

Dijkstra's VS A*


- Dijkstra's algorithm visits in all directions



- A* visits mainly towards the goal, but does not hedge its bets to ensure optimality




Dijkstra's VS A*



Reading Resources

- If you are really interested in planning
- LaValle, Steven M. Planning algorithms. Cambridge university press, 2006.



Next Lecture

- Trajectory planning
 - Guest Lecture by Haokun Wang
 - Aerial Robot (the famous one in HKUST Robotics Institute) 😊

