

My Lecture Notes

Overall Robot Navigation Paradigm

External World-> Sense -> Estimate -> Plan -> Act -> External World (Repeat)

- Sense+Estimate: estimate current and past pose
- Planning: Generate future robot pose
- Control: Stabilize robot pose

Main Questions to Solve

1. Where am I? (Sensing/Estimation)
2. Where am I going? (Planning)
3. How do I get there? (Control)

L2: Pose and Rotation

- Pose = position (translation) + orientation (rotation)
- Rigid Body: no deformation
 - Displacement:
 - * Lengths preserved
 - * Cross product preserved
 - Motion
 - * include both translation and rotation
- Rotation
 - Right-handed coordinates
 - Can be represented with rotation matrices, Euler angles, Quaternions, etc.
 - * Rotation matrix have no singularity and redundant parameters
 - * Euler angles have gimble lock
 - * Quaternion solves above problem

L3: Localization and Kinematics

- Solutions to “Where am I?” problem
 - external infrastructures (eg. GPS)
 - Odometry
 - * Pose estimation only
 - * Pros:
 - low computation complexity
 - high freq
 - * Cons:
 - drift
 - no map
 - SLAM
 - * Pose + local mapping

- * Pros:
 - builds local map
 - low drift
 - * Cons:
 - higher complexity
 - map redundancy for long-term use
 - no global localization
- Map-based Localization
 - * Pose estimation
 - * Pros
 - low drift with good map
 - lower complexity
 - good for long-term use
 - global localization
 - * Cons
 - need good map
- Wheel Odometry
 - Locomotion: ability to transport/move itself place to place
 - * issues: stability, contact, environment
 - wheels
 - * standard wheel (2 dof (degree of freedom))
 - * castor wheel (2 dof)
 - * swedish wheel (3 dof)
 - * ball wheel
 - estimate with kinematics (study of motion without regard to forces)

L4: Sensors

- Interoceptive: within body
 - Accelerometer
 - * measures translational acceleration
 - Gyroscope
 - * measures angular rate
 - IMU
 - * measures linear acceleration and angular velocity
 - * Pros
 - available and outlier free
 - high rate measurements
 - low cost
 - * Cons
 - noisy
 - can't operate when inertial and visual measurements not in same frame (eg. on train)
- Exteroceptive Sensors
 - GNSS (satellite)
 - * GPS

- needs at least 4 satellites to estimate position
 - doesn't work well in denied areas (urban, forest, park, tunnel, etc)
 - Camera
 - * Monocular
 - simplest
 - unknown depth
 - * Uses perspective projection
 - closer object appear bigger
 - * For solving depth
 - can use pictorial cues
 - stereo camera can compute depth
 - Lidars
 - * Pros
 - wide fov
 - precise depth
 - illumination free
 - * Cons
 - no RGB info
 - higher cost
 - sparse points per scan
 - Radar
 - * Pros
 - weather robust
 - speed
 - * Cons
 - noisy
 - less support
 - RGB-D Sensor
 - * Lidar + Monocular
 - * Pros
 - fast speed
 - relatively good depth
 - * Cons
 - does not work outdoor
 - affected by illumination
- Multi-sensor Fusion
 - combines sensors
 - Pros
 - * more robust
 - * higher precision
 - Cons
 - * high cost
 - * time sync

L5: ICP

- 3D pose estimation with known correspondences
 - compute SVD
 - get rotation and translation, compute pose
- For unknown correspondences, can use ICP
- ICP Idea: iteratively find alignment
- Algo:
 1. start with some guess
 2. find nearest neighbor
 3. refine rotation and translation based on latest data
 4. repeat step 2 until convergence
- Sampling/Filters
 - Can have uniform, random, feature-based sampling, etc.
- Weighting Correspondences
 - Noise: weight based on uncertainty
 - Outlier: lower weights for higher point-to-point distances
 - Determine transformation that minimizes weighted error function
- Data Association
 - closest point, point-2-plane, etc.

L6: Map

- Representations
 - Point Cloud
 - * Pros
 - no discretization of data
 - map area not limited
 - * Cons
 - unbounded memory usage
 - no direct representation of free or unknown space
 - Grid
 - * Pros
 - volumetric representation
 - constant access time
 - friendly to planning
 - * Cons
 - memory requirement
 - extent of map has to be known/guessed
 - Octree
 - * tree-based data structure
 - * Pros
 - full 3D model
 - memory efficient
 - * Cons
 - tricky to implement

- Elevation
 - * 2D grid that stores estimated height for each cell
 - * Pros
 - 2.5D representation
 - constant time access
 - * Cons
 - no vertical objects
- Features
 - * Use features such as landmarks, objects, etc.
 - * Pros
 - less memory loss
 - support localization
 - close to human eye perception
 - * Cons
 - can't represent complex env
 - not planning oriented

L7: Bayes Filter

- Problem Overview:
 - Move (increases uncertainty)
 - Predict (use sensors, inputs, etc.)
 - Sense (get sensor measurements)
 - Update (decrease uncertainty)
- Bayes filter problem example:
 - Suppose robot obtains measurement z (eg. brightness)
 - What is $P(\text{door open}|z)$

L8: Particle Filters

- Particle filters are implementations of recursive Bayesian filtering
- Particles are propagated according to the motion model
- weighted according to likelihood of observations
- Pros
 - Easy to implement
 - able to handle nonlinear systems without linearization
 - able to represent arbitrary distribution
- Cons
 - particle degeneracy problem
 - need lots of particles to represent high dimensional state space, more computation complexity as well
- Applications
 - robot pose tracking, target tracking, etc.

L9: Kalman Filter and EKF

- Particle Degeneneracy Problem:
 - situation where a small number of particles in the filter carry a significant weight
 - eg. filter out 25% particles each time, after 4 operations particles originate from same source
- To solve problem, can use Kalman Filter (KF) and Extended Kalman Filter (EKF)
- Kalman Filter
 - weighted mean with Gaussians
 - Steps:
 1. Initialize system state
 2. Predict next step state
 3. Update with Kalman Gain
 4. Make correction
 5. Return to predict with new mean and variance
 - Everything stays Gaussian
 - variance never increases
 - prediction and update can happen in arbitrary order
 - Pros
 - * simple
 - * efficient
 - Cons
 - * assumes linearity and Gaussian
- Kalman Gain
 - Intuition: How much to trust the sensor vs prediction
 - example: $R=0$ (perfect sensor)
 - * $\text{gain} = C^{-1}$
 - * variance is 0
 - example: $R=\infty$
 - * gain is 0
- EKF
 - Resolves non-linear functions by using local linearization
 - generates Gaussian approximation

L10: EKF SLAM

- Overall Goal: obtain both feature map and robot poses in real time
- SLAM Problem:
 - Given
 - * control commands (odometry)
 - * observations (measurements)
 - Want
 - * Map of env
 - * Path or current pose of robot

- Difficulties
 - * real world mapping between observations and landmarks unknown
 - * picking wrong data associations can have large consequences
 - * chicken-or-egg problem
 - a map is needed for localization
 - pose estimate needed for mapping
- EKF SLAM
 - Steps:
 - * State prediction
 - state propagation with motion model
 - * Measurement prediction: predict where map landmark should be if it appears
 - * Obtained measurement: observe landmark with sensor
 - * Difference: get difference between predicted and observed
 - * Update step: update map and pose
 - Convergence results for linear case, diverge if nonlinearities large

L11: Place Recognition

- Loop Closing
 - recognizing already mapped area
 - uncertainties collapse
- Lidar Place Recognition
 - check if point cloud has been seen before
 - need retrieval scheme
 - * eg. ring key
- Visual Place Recognition
 - check if image were taken around same location
 - difficult because challenges like information-rich, illumination changes, etc.

L12: Pose Graph SLAM

- Idea of Graph SLAM
 - use graph to represent problem
 - every node corresponds to pose during mapping
 - every edge is spatial constraint
 - build graph and find node config that minimizes error introduced by constraints

L13: Graph SLAM with Landmarks

- Rank of matrix H
 - 2D landmark-pose constraint
 - * blocks J_{ij} are 2x3 matrices

- * H_{ij} cannot have more than rank 2
 - bearing-only constraint
 - * J_{ij} are 1x3 matrices
 - * H_{ij} has rank 1
- Questions
 - How many 2D landmark observations are needed to resolve robot pose?
 - * Need min 2
 - How many bearing-only observations are needed to resolve robot pose?
 - * Need min 3

L14: Visual Feature Detection

- Advantages of local features
 - locality
 - distinctiveness
 - quantity
 - efficiency
 - generality
- feature detection summary
 - compute gradient at each point in image
 - compute H matrix from gradient entries
 - compute eigenvalues
 - find points with large response
 - choose points where lambda is local maximum as features

L15: Visual Descriptor and Matching

- Transformational invariance
 - Want features to be same regardless of transformation
 - Need to have both of following
 1. Detector is invariant
 2. Design invariant feature descriptor
- Scale Invariant Feature Transform (SIFT)
 - basic idea
 - * take 16x16 square window around feature
 - * compute edge orientation for each pixel
 - * throw out weak edges
 - * create histogram of surviving edge orientations
 - properties
 - * translation invariant
 - * illumination invariant
 - * rotation invariant

- Random Sample Consensus (RANSAC)
 - used to remove outliers from data
 - uses trial error approach to find best model

L16: Planning and Graph Construction

- motion planning (or piano mover's problem) finds optimal route to goal
- configuration
 - configuration is specification of every point on robot body
 - configuration space (c-space) is all possible robot configurations
 - use c-space instead of workspace for motion planning as robot is represented as point (easier for motion planning)
 - representing obstacle in c-space can be complicated, so use approximation
 - c-space is discretized for path planning
- two approaches to discretize c-spaces:
 - combinatorial planning
 - sampling-based planning
- combinatorial planning
 - visibility graph
 - * construct path as polygonal line through vertices of obstacles
 - * pros: simple
 - * cons: path too close to obstacles (not safe)
 - voronoi diagram
 - * take midpoint between obstacles to maximize distance between robot and obstacles
 - * pros: safe, good clearance for uncertain robot
 - * cons: complex calculation, not suitable for short-distance sensors
 - approximate cell decomposition
 - * decompose into cells, construct connectivity graph, then search
 - * pros: simple computations, simpler
 - * cons: large storage space
- sampling based planning
 - probabilistic roadmap (PRM)
 - * build graph to characterize free config space, use graph search to find path
 - * pros: can cope with high dimensions, probabilistically complete (will contain a solution if it exists)
 - * cons: collision detection takes majority of time, suboptimal if only limited samples given, need whole c-space as prior condition

L17: RTT, A-Star, Dijkstra

- rapidly exploring random trees (RTT)

- aggressively probe and explore c-space by expanding incrementally from initial config
 - pros: balance between greedy search and exploration, easy to implement
 - cons: metric sensitivity, unknown rate of convergence
- Dijkstra's Algo
 - pros: complete and optimal
 - cons: can only see cost accumulated so far, no information about goal location
- A-Star
 - combines dijkstra and a heuristic